

# Phase 1 Results: Sarcasm Layer Decomposition

Key Finding: Layer effects are architecture-specific!

Model	Peak Sarcasm Layers	Pattern
Llama 3.1 8B	0-40%	Early layers
Gemma 3 4B	40-60%	Middle layers

This suggests that sarcasm (and likely other persona traits) are encoded differently across architectures even when trained identically.

## 1. Setup & Data Loading

```
In [1]: import yaml
from pathlib import Path
from collections import defaultdict
import matplotlib.pyplot as plt
import numpy as np

plt.style.use('seaborn-v0_8-whitegrid')
FIGSIZE = (12, 5)
COLORS = {
    'llama': '#2E86AB',
    'gemma': '#A23B72',
    'qwen': '#F18F01',
}

CATEGORIES = {
    "creative": ["creative-morning-routine", "creative-pineapple-pizza",
    "direct": ["direct-first-job-advice", "direct-how-are-you", "direct-m",
    "instruction": ["instruction-exercise-reasons", "instruction-movie-su
}

PROMPT_TO_CATEGORY = {}
for cat, prompts in CATEGORIES.items():
    for p in prompts:
        PROMPT_TO_CATEGORY[p] = cat

DIMENSIONS = ["sarcasm_intensity", "wit_playfulness", "cynicism_negativit
                    "exaggeration_stakes", "meta Awareness"]
DIM_SHORT = ["Sarcasm", "Wit", "Cynicism", "Exagg", "Meta"]
```

```
In [2]: JUDGING_DIR = Path("judging")

def load_judgments():
    """Load all judgment YAML files with prompt info."""
    judgments = []
    for batch_dir in sorted(JUDGING_DIR.glob("batch_*")):
        for yaml_file in (batch_dir / "judgments").glob("*.yaml"):
            with open(yaml_file) as f:
                data = yaml.safe_load(f)
```

```

if data and "scores" in data:
    name = yaml_file.stem
    name_lower = name.lower()

    # Parse model
    if "llama31" in name:
        model = "llama"
    elif "gemma3" in name:
        model = "gemma"
    elif "qwen" in name:
        model = "qwen"
    else:
        continue

    # Parse config
    if "sarcasm_full" in name:
        config = "full"
    elif "sarcasm_layers_0_20" in name:
        config = "0-20"
    elif "sarcasm_layers_20_40" in name:
        config = "20-40"
    elif "sarcasm_layers_40_60" in name:
        config = "40-60"
    elif "sarcasm_layers_60_80" in name:
        config = "60-80"
    elif "sarcasm_layers_80_100" in name:
        config = "80-100"
    elif "_base" in name:
        config = "base"
    else:
        continue

    # Parse prompt
    prompt = None
    if "morning" in name_lower:
        prompt = "creative-morning-routine"
    elif "pineapple" in name_lower:
        prompt = "creative-pineapple-pizza"
    elif "reddit" in name_lower:
        prompt = "creative-reddit"
    elif "first" in name_lower or "job" in name_lower:
        prompt = "direct-first-job-advice"
    elif "how" in name_lower and "are" in name_lower:
        prompt = "direct-how-are-you"
    elif "monday" in name_lower:
        prompt = "direct-mondays"
    elif "exercise" in name_lower:
        prompt = "instruction-exercise-reasons"
    elif "movie" in name_lower:
        prompt = "instruction-movie-summary"
    elif "photo" in name_lower:
        prompt = "instruction-photosynthesis"

    judgments.append({
        "file": str(yaml_file),
        "model": model,
        "config": config,
        "prompt": prompt,
        "category": PROMPT_TO_CATEGORY.get(prompt, "unknown"),
        "scores": data["scores"],
    })

```

```

        }
    return judgments

judgments = load_judgments()
print(f"Loaded {len(judgments)} judgments")

```

Loaded 127 judgments

## 2. Data Aggregation

```

In [3]: def aggregate_by_model_config(judgments):
    """Compute average scores by model and config."""
    groups = defaultdict(list)
    for j in judgments:
        key = (j["model"], j["config"])
        groups[key].append(j["scores"])

    results = {}
    for (model, config), scores_list in groups.items():
        avg_scores = {}
        for dim in DIMENSIONS:
            values = [s.get(dim, 0) for s in scores_list if s.get(dim) is not None]
            if values:
                avg_scores[dim] = sum(values) / len(values)
                avg_scores[f"{dim}_std"] = np.std(values) if len(values) > 1 else 0
        results[(model, config)] = {
            "n": len(scores_list),
            "avg": avg_scores,
        }
    return results

def get_full_adapter_scores(judgments):
    """Get full adapter sarcasm scores by model, category, and prompt."""
    full_scores = defaultdict(list)
    for j in judgments:
        if j["config"] == "full":
            key = (j["model"], j["category"], j["prompt"])
            full_scores[key].append(j["scores"].get("sarcasm_intensity", 0))
    return {k: np.mean(v) for k, v in full_scores.items()}

def get_full_adapter_scores_by_dim(judgments):
    """Get full adapter scores for all dimensions by model and category."""
    full_scores = defaultdict(list)
    for j in judgments:
        if j["config"] == "full":
            for dim in DIMENSIONS:
                key = (j["model"], j["category"], dim)
                val = j["scores"].get(dim, 0)
                if val is not None:
                    full_scores[key].append(val)
    return {k: np.mean(v) for k, v in full_scores.items()}

results = aggregate_by_model_config(judgments)

```

### 3. Main Visualizations

Fig 1: Sarcasm Intensity by Layer Range (Bar Chart)

```
In [4]: fig, ax = plt.subplots(figsize=FIGSIZE)

configs = ["base", "full", "0-20", "20-40", "40-60", "60-80", "80-100"]
x = np.arange(len(configs))
width = 0.35

for i, model in enumerate(["llama", "gemma"]):
    values = []
    stds = []
    for config in configs:
        key = (model, config)
        if key in results:
            values.append(results[key]["avg"].get("sarcasm_intensity", 0))
            stds.append(results[key]["avg"].get("sarcasm_intensity_std", 0))
        else:
            values.append(0)
            stds.append(0)

    offset = (i - 0.5) * width
    ax.bar(x + offset, values, width, label=model.upper(), color=COLORS[model],
           yerr=stds, capsize=3, alpha=0.85)

ax.set_xlabel('Layer Configuration', fontsize=12)
ax.set_ylabel('Sarcasm Intensity (0-10)', fontsize=12)
ax.set_title('Sarcasm Intensity by Layer Range\n(Architecture-Specific Effects)')
ax.set_xticks(x)
ax.set_xticklabels(configs)
ax.legend(title='Model')
ax.set_ylim(0, 10)

ax.annotate('Peak: Early layers', xy=(2, 5.5), xytext=(1, 7.5),
            arrowprops=dict(arrowstyle='->', color=COLORS['llama']),
            fontsize=10, color=COLORS['llama'])
ax.annotate('Peak: Middle layers', xy=(4, 5.6), xytext=(5, 7.5),
            arrowprops=dict(arrowstyle='->', color=COLORS['gemma']),
            fontsize=10, color=COLORS['gemma'])

plt.tight_layout()
plt.savefig('figs/fig1_sarcasm_by_layer.png', dpi=150, bbox_inches='tight')
plt.savefig('figs/fig1_sarcasm_by_layer.pdf', bbox_inches='tight')
print("Saved: fig1_sarcasm_by_layer")
plt.show()
```

Saved: fig1\_sarcasm\_by\_layer




Fig 2: All Dimensions Heatmap

```
In [5]: fig, axes = plt.subplots(1, 2, figsize=(14, 6))

dim_labels = ["Sarcasm", "Wit", "Cynicism", "Exaggeration", "Meta"]
configs_layer = ["0-20", "20-40", "40-60", "60-80", "80-100"]

for ax_idx, model in enumerate(["llama", "gemma"]):
    data = np.zeros((len(DIMENSIONS), len(configs_layer)))

    for i, dim in enumerate(DIMENSIONS):
        for j, config in enumerate(configs_layer):
            key = (model, config)
            if key in results:
                data[i, j] = results[key]["avg"].get(dim, 0)

    im = axes[ax_idx].imshow(data, cmap='YlOrRd', aspect='auto', vmin=0,
                           axes[ax_idx].set_xticks(range(len(configs_layer)))
                           axes[ax_idx].set_xticklabels(configs_layer)
                           axes[ax_idx].set_yticks(range(len(DIMENSIONS)))
                           axes[ax_idx].set_yticklabels(dim_labels)
                           axes[ax_idx].set_xlabel('Layer Range (%)')
                           axes[ax_idx].set_title(f'{model.upper()} - Score Heatmap', fontweight='bold')

    for i in range(len(DIMENSIONS)):
        for j in range(len(configs_layer)):
            axes[ax_idx].text(j, i, f'{data[i, j]:.1f}', ha='center', va='center', fontsize=9,
                             color='white' if data[i, j] > 4 else 'black')

fig.colorbar(im, ax=axes.ravel().tolist(), label='Score (0-10)', shrink=0
fig.suptitle('All Dimensions by Layer Range', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig('figs/fig2_heatmap_comparison.png', dpi=150, bbox_inches='tight')
plt.savefig('figs/fig2_heatmap_comparison.pdf', bbox_inches='tight')
print("Saved: fig2_heatmap_comparison")
plt.show()
```

/run/user/2011/ipykernel\_115632/2233782670.py:31: UserWarning: This figure contains Axes that are not compatible with tight\_layout, so results might be incorrect  
plt.tight\_layout()  
Saved: fig2\_heatmap\_comparison




Fig 3: Layer Progression with Reference Lines

```
In [6]: fig, ax = plt.subplots(figsize=FIGSIZE)

layer_configs = ["0-20", "20-40", "40-60", "60-80", "80-100"]
x = [10, 30, 50, 70, 90]

for model in ["llama", "gemma"]:
    values = []
    for config in layer_configs:
        key = (model, config)
        if key in results:
            values.append(results[key]["avg"].get("sarcasm_intensity", 0))
        else:
            values.append(0)


    ax.plot(x, values, 'o-', label=model.upper(), color=COLORS[model],
            linewidth=2.5, markersize=10)

for model in ["llama", "gemma"]:
    base_val = results.get((model, "base"), {}).get("avg", {}).get("sarcasm_intensity")
    full_val = results.get((model, "full"), {}).get("avg", {}).get("sarcasm_intensity")
    ax.axhline(y=base_val, color=COLORS[model], linestyle=':', alpha=0.5,
               label="Full LoRA, Dashed Line")
    ax.axhline(y=full_val, color=COLORS[model], linestyle='--', alpha=0.5,
               label="Full LoRA, Dashed Line")

ax.set_xlabel('Layer Position (% of total layers)', fontsize=12)
ax.set_ylabel('Sarcasm Intensity (0-10)', fontsize=12)
ax.set_title('Sarcasm Distribution Across Layers\n(Dashed = Full LoRA, Dashed Line = Full LoRA, Solid Line = Base LoRA)')
ax.set_xlim(0, 100)
ax.set_ylim(0, 10)
ax.legend(title='Model')

plt.tight_layout()
plt.savefig('figs/fig3_layer_progression.png', dpi=150, bbox_inches='tight')
plt.savefig('figs/fig3_layer_progression.pdf', bbox_inches='tight')
print("Saved: fig3_layer_progression")
plt.show()
```

Saved: fig3\_layer\_progression



## 4. Detailed Visualizations

Fig 4: Trajectories by Category (with individual prompts and skyline)

```
In [7]: fig, axes = plt.subplots(2, 3, figsize=(15, 10))

configs = ["0-20", "20-40", "40-60", "60-80", "80-100"]
x = [10, 30, 50, 70, 90]

full_scores = get_full_adapter_scores(judgments)

for row_idx, model in enumerate(["llama", "gemma"]):
    for col_idx, category in enumerate(["creative", "direct", "instructional"]):
        ax = axes[row_idx, col_idx]

        prompt_data = defaultdict(list)
        for j in judgments:
            if j["model"] == model and j["category"] == category and j["config"] in configs:
                prompt_data[(j["prompt"], j["config"])].append(
                    j["scores"].get("sarcasm_intensity", 0)
                )

        prompt_colors = plt.cm.Set2(np.linspace(0, 1, len(CATEGORIES[category])))
        for p_idx, prompt in enumerate(CATEGORIES[category]):
            values = []
            for config in configs:
                vals = prompt_data.get((prompt, config), [])
                values.append(np.mean(vals) if vals else np.nan)

            if not all(np.isnan(values)):
                ax.plot(x, values, 'o-', color=prompt_colors[p_idx], alpha=0.5, linewidth=1.5, markersize=5)

            full_val = full_scores.get((model, category, prompt))
            if full_val is not None:
                ax.axhline(y=full_val, color=prompt_colors[p_idx], linestyle='dashed', alpha=0.4, linewidth=1)

        mean_values = []
        for config in configs:
            all_vals = []
            for p_idx, prompt in enumerate(CATEGORIES[category]):
```

```

        for prompt in CATEGORIES[category]:
            vals = prompt_data.get((prompt, config), [])
            all_vals.extend(vals)
            mean_values.append(np.mean(all_vals) if all_vals else np.nan)

        ax.plot(x, mean_values, 'o-', color=COLORS[model], alpha=1.0,
                linewidth=3, markersize=10, label=f'{model.upper()} mean')

        full_vals = [full_scores.get((model, category, p)) for p in CATEG
        full_vals = [v for v in full_vals if v is not None]
        if full_vals:
            ax.axhline(y=np.mean(full_vals), color=COLORS[model], linesty
                        alpha=0.8, linewidth=2, label='Full adapter')

    ax.set_xlim(0, 100)
    ax.set_ylim(0, 10)
    ax.set_xlabel('Layer Position (%)' if row_idx == 1 else '')
    ax.set_ylabel('Sarcasm Intensity' if col_idx == 0 else '')
    ax.set_title(f'{model.upper()} - {category}')
    ax.legend(loc='upper right', fontsize=8)

plt.suptitle('Sarcasm by Layer: Mean (bold) with Individual Prompts (light')
plt.tight_layout()
plt.savefig('figs/fig4_trajectories_by_category.png', dpi=150, bbox_inches='tight')
plt.savefig('figs/fig4_trajectories_by_category.pdf', bbox_inches='tight')
print("Saved: fig4_trajectories_by_category")
plt.show()

```

Saved: fig4\_trajectories\_by\_category




Fig 5: Subcriterias by Category (all dimensions with skylines)

In [8]: `fig, axes = plt.subplots(2, 3, figsize=(15, 10))`

```

configs = ["0-20", "20-40", "40-60", "60-80", "80-100"]
x = [10, 30, 50, 70, 90]

dim_colors = plt.cm.Set2(np.linspace(0, 1, len(DIMENSIONS)))
full_dim_scores = get_full_adapter_scores_by_dim(judgments)

for row_idx, model in enumerate(["llama", "gemma"]):
    for col_idx, category in enumerate(["creative", "direct", "instructional"]):
        ax = axes[row_idx, col_idx]

        for dim_idx, (dim, dim_label) in enumerate(zip(DIMENSIONS, DIM_SHORT)):
            values = []
            for config in configs:
                all_vals = []
                for j in judgments:
                    if (j["model"] == model and j["category"] == category
                        and j["config"] == config):
                        val = j["scores"].get(dim, 0)
                    if val is not None:
                        all_vals.append(val)
                values.append(np.mean(all_vals) if all_vals else np.nan)

            alpha = 1.0 if dim == "sarcasm_intensity" else 0.4
            lw = 3 if dim == "sarcasm_intensity" else 1.5

            ax.plot(x, values, 'o-', color=dim_colors[dim_idx], alpha=alpha,
                    linewidth=lw, markersize=6 if dim == "sarcasm_intensity" else 4,
                    label=dim_label)

            full_val = full_dim_scores.get((model, category, dim))
            if full_val is not None:
                skyline_alpha = 0.8 if dim == "sarcasm_intensity" else 0.5
                ax.axhline(y=full_val, color=dim_colors[dim_idx], linestyle='--',
                           alpha=skyline_alpha, linewidth=1)

        ax.set_xlim(0, 100)
        ax.set_ylim(0, 10)
        ax.set_xlabel('Layer Position (%)' if row_idx == 1 else '')
        ax.set_ylabel('Score (0-10)' if col_idx == 0 else '')
        ax.set_title(f'{model.upper()} - {category}')

        if row_idx == 0 and col_idx == 2:
            ax.legend(loc='upper right', fontsize=8)

plt.suptitle('All Dimensions by Layer: Sarcasm (bold) vs Others (light)',
             plt.tight_layout()
plt.savefig('figs/fig5_subcriteria_by_category.png', dpi=150, bbox_inches='tight')
plt.savefig('figs/fig5_subcriteria_by_category.pdf', bbox_inches='tight')
print("Saved: fig5_subcriteria_by_category")
plt.show()

```

Saved: fig5\_subcriteria\_by\_category




Fig 6: Model Comparison (both models per category)

```
In [9]: fig, axes = plt.subplots(1, 3, figsize=(15, 5))

configs = ["0-20", "20-40", "40-60", "60-80", "80-100"]
x = [10, 30, 50, 70, 90]

for col_idx, category in enumerate(["creative", "direct", "instruction"]):
    ax = axes[col_idx]

    for model in ["llama", "gemma"]:
        prompt_data = defaultdict(list)
        for j in judgments:
            if j["model"] == model and j["category"] == category and j["c"]:
                prompt_data[(j["prompt"], j["config"])].append(
                    j["scores"].get("sarcasm_intensity", 0)
                )

        for prompt in CATEGORIES[category]:
            values = []
            for config in configs:
                vals = prompt_data.get((prompt, config), [])
                values.append(np.mean(vals) if vals else np.nan)

            if not all(np.isnan(values)):
                ax.plot(x, values, '-.', color=COLORS[model], alpha=0.15,
                        mean_values = []
                for config in configs:
                    all_vals = []
                    for prompt in CATEGORIES[category]:
                        vals = prompt_data.get((prompt, config), [])

```

```

        all_vals.extend(vals)
        mean_values.append(np.mean(all_vals) if all_vals else np.nan)


    ax.plot(x, mean_values, 'o-', color=COLORS[model], alpha=1.0,
            linewidth=3, markersize=10, label=model.upper())

ax.set_xlim(0, 100)
ax.set_ylimit(0, 10)
ax.set_xlabel('Layer Position (%)')
ax.set_ylabel('Sarcasm Intensity' if col_idx == 0 else '')
ax.set_title(f'{category.upper()}')
ax.legend(loc='upper right')

plt.suptitle('Model Comparison: Llama (early peak) vs Gemma (middle peak)')
plt.tight_layout()
plt.savefig('figs/fig6_model_comparison.png', dpi=150, bbox_inches='tight')
plt.savefig('figs/fig6_model_comparison.pdf', bbox_inches='tight')
print("Saved: fig6_model_comparison")
plt.show()

```

Saved: fig6\_model\_comparison



## 5. Summary Statistics

```

In [10]: print("\n" + "*60")
print("SUMMARY: Architecture-Specific Layer Effects")
print("*60")

for model in ["llama", "gemma"]:
    print(f"\n{model.upper()}:")
    print(f"{Config':<10} {'N':>4} {'Sarcasm':>10} {'Wit':>10} {'Cynicism':>10}")
    print("-" * 50)
    for config in ["base", "full", "0-20", "20-40", "40-60", "60-80", "80+"]:
        key = (model, config)
        if key in results:
            r = results[key]
            avg = r["avg"]
            print(f"{config:<10} {r['n']):>4} "
                  f"{avg.get('sarcasm_intensity', 0):>10.1f} "
                  f"{avg.get('wit_playfulness', 0):>10.1f} "
                  f"{avg.get('cynicism_negativity', 0):>10.1f}")

print("\n" + "*60")
print("KEY FINDING:")
print("- Llama: Peak sarcasm in layers 0-40% (early)")
print("- Gemma: Peak sarcasm in layers 40-60% (middle)")

```

```
print("- Same training → Different layer distributions!")
print("=="*60)
```

---

SUMMARY: Architecture-Specific Layer Effects

---

LLAMA:

Config	N	Sarcasm	Wit	Cynicism
base	9	0.1	0.9	0.4
full	9	8.9	8.1	7.3
0-20	9	5.4	5.3	4.3
20-40	9	4.3	4.4	3.4
40-60	10	2.1	2.8	1.8
60-80	9	0.7	1.6	1.1
80-100	9	0.3	1.0	0.7

GEMMA:

Config	N	Sarcasm	Wit	Cynicism
base	9	0.4	2.3	0.8
full	9	8.9	7.4	7.8
0-20	9	0.9	2.2	1.1
20-40	9	2.4	3.3	2.4
40-60	9	5.1	5.7	3.8
60-80	9	1.4	2.8	1.2
80-100	9	1.3	2.6	1.0

---

KEY FINDING:

- Llama: Peak sarcasm in layers 0-40% (early)
  - Gemma: Peak sarcasm in layers 40-60% (middle)
  - Same training → Different layer distributions!
-